

## 15-859E Assignment 2: Multigrid

Franklin Chen

October 21, 1998

## Contents

<b>1</b>	<b>Equation</b>	<b>3</b>
1.1	Source . . . . .	3
1.2	Boundary values . . . . .	3
<b>2</b>	<b>Implementation language</b>	<b>4</b>
<b>3</b>	<b>Platform</b>	<b>5</b>
<b>4</b>	<b>Algorithms</b>	<b>6</b>
4.1	Relaxation . . . . .	6
4.2	Interpolation . . . . .	6
4.3	Restriction . . . . .	6
4.4	Multigrid . . . . .	6
<b>5</b>	<b>Command line arguments</b>	<b>7</b>
<b>6</b>	<b>Timing, accuracy, speed</b>	<b>8</b>
<b>7</b>	<b>Visual presentations</b>	<b>8</b>
7.1	Plot . . . . .	8
7.2	Picture . . . . .	8

## List of Figures

1	Iterative method performances . . . . .	9
---	---	---

## 1 Equation

I used the *Poisson* equation in 2d:

$$\nabla^2 \mathbf{u} = \mathbf{f}$$

This was suggested in the assignment, and was also discussed in Briggs.

### 1.1 Source

I tried some sample  $\mathbf{f}$  for the equation. One was  $\mathbf{f} = \mathbf{0}$  (Laplace). I settled on another constant, just for the heck of it.

I did not provide a user interface for plugging in various  $\mathbf{f}$ , though that would be easy, e.g., parse some stream of numbers from a file.

### 1.2 Boundary values

The *Dirichlet boundary values* I chose were to take the 2d unit square as the boundaries, and have all sides be 0 except a portion of sine wave on one side, in order to get a simple but not completely trivial solution surface.

## 2 Implementation language

I implemented everything in *Objective Caml* (O'Caml), a dialect of ML. I hadn't done much with this dialect, but figured now was as good a time as any to experiment with it, instead of using *Standard ML*.

Since O'Caml supports arrays of unboxed floating point numbers, this gave a fighting chance of not being too inefficient, though I suspect that I could write C or C++ code that does significantly better.

Because I lacked time, I did not experiment with using a language with high level support of arrays, such as ZPL, SAC, or FiSH. I am interested in trying this at some point.

Note: I use arrays of arrays for matrices. This obviously misses out on the memory locality and opportunities for good caching and loop unrolling possible when using a block of contiguous memory as a matrix (as the `svl` library in C++ attempts to do, as does the SML/NJ library for Standard ML).

My first implementation was written purely functionally, with no side effects in the computation, e.g., arrays were never modified. I reimplemented it to use mutation, and the speedup wasn't all that large, actually. I stuck to the mutative version though.

My primary source of equations was Briggs. I also looked at *Numerical Recipes in C* earlier, did not consult it while in the process of writing my own, since I first wrote all my code purely functionally, then converted it in steps to use mutation.

### **3 Platform**

I developed my program on my home PC, a Pentium 200 with 32 MB RAM running Red Hat Linux 5.1. I did some runs on my office PC, a Pentium II 400 with 128 MB RAM, and performance appeared to be around four times greater.

## 4 Algorithms

### 4.1 Relaxation

I implemented both *Gauss-Seidel* (red-black) and *Jacobi*, both of them with optional weighting ( $\omega$ ) for overrelaxation.

I made no effort to make these crucial steps fast. A lot of interesting cache-related optimizations could be performed, but I could only do that if I were using something as low level as C.

### 4.2 Interpolation

I used 2d linear interpolation, as suggested by Briggs.

### 4.3 Restriction

I used full-weighted 2d restriction, as suggested by Briggs.

### 4.4 Multigrid

I implemented the  $\mu$ -cycle generalization of the V-cycle, with fixed numbers of iterations, specified for the duration of a run from the command line.

## 5 Command line arguments

Many parameters can be set for the duration of a run by means of command line arguments.

Usage:

```
-escape Escape from algorithm when relative residual small?
    Default = false
-debug Debugging level
    Default = 0
    0 = no debug
    1 = output initial and final states
    2 = output after every relaxation step
    3 = output after each read and black step of Gauss-Seidel
    4 = output full multigrid residual and its restriction
-nu0 Number of mu-cycles for each phase of full multigrid
    Default = 1
-nu1 Number of pre-cycle relaxation steps in mu-cycle
    Default = 1
-nu2 Number of post-cycle relaxation steps in mu-cycle
    Default = 1
-mu Number of recursive calls in a mu-cycle
    Default = 1
-coarse Size of coarsest grid base case ( $2^k+1$ )
    Default = 3
-size Size of finest grid base case ( $2^k+1$ )
    Default = 9
-omega Weight for relaxation method (in  $(0, 1]$ )
    Default = 0.6666666666667
-relax Relaxation method (gs or jacobi)
    Default = gs
    gs = Gauss-Seidel
    jacobi = Jacobi
-multigrid Overall method
    Default = full
    full = full multigrid
    mucycle = mu-cycle
    gs = Gauss-Seidel
    jacobi = Jacobi

-iterations Number of iterations of relaxation
    Default = 1000
    (Only applicable to -multigrid gs|jacobi)
```

## 6 Timing, accuracy, speed

Roughly, I was able to get full multigrid running to the desired accuracy suggested (relative residual  $< 10^{-6}$ ) on a  $256 \times 256$  grid in about 9 seconds on my home PC, and about 2 seconds on my office PC.

Attempts to get fine-grained data for algorithms other than full multigrid were not very successful, because for larger  $n$  things would take forever, and furthermore, not achieve good accuracy.

I generated a bunch of runs by trying to guess parameters. I tried to be more scientific by adding code to detect when the relative residual was small enough, but this expensive computation slowed things down significantly. I didn't try to just periodically do the check instead.

## 7 Visual presentations

### 7.1 Plot

In Figure 1, we show the results for

- Weighted red-black Gauss-Seidel, 1300 iterations (note last two grid sizes did not complete in a reasonable time).
- $\mu$ -cycle, actually just a V-cycle, with 500 iterations pre and post.
- Full multigrid, with parameters 4, 2, 2.

### 7.2 Picture

I didn't have time to do anything fancy with graphics, an area I don't really know much about (but wish to). My program will generate a simplistic VRML file, however, treating the solution field as an `ElevationGrid`. An example is in `dataset.data` (see `Makefile` for options used).



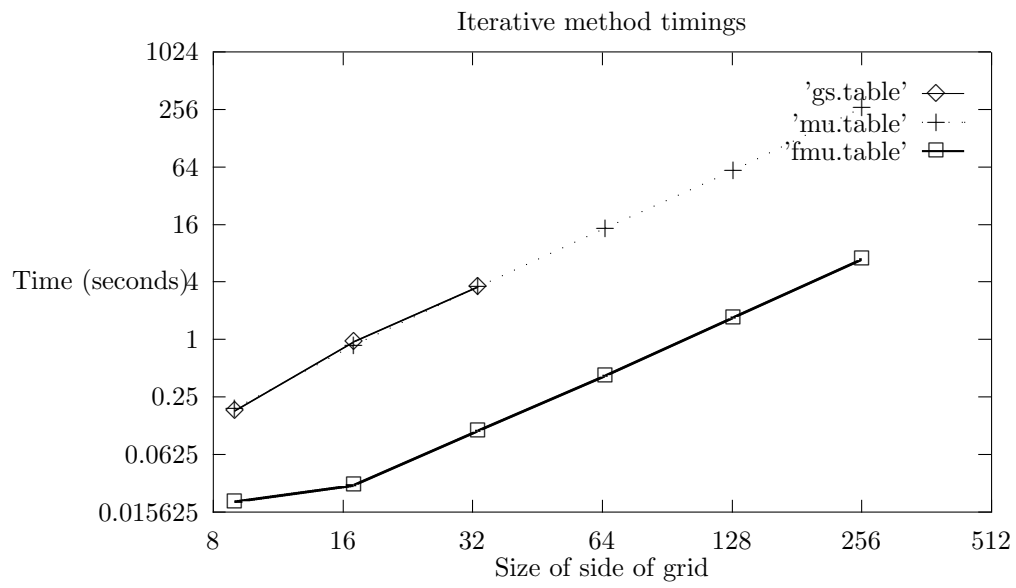


Figure 1: Iterative method performances